

Introduction to dplyr

Paul M. Magwene

What is dplyr?

dplyr is a package that provides a “grammar for data manipulation”

Key “verbs” in the dplyr package:

- ▶ `select()`
- ▶ `filter()`
- ▶ `mutate()`
- ▶ `arrange()`
- ▶ `summarize()`
- ▶ `group_by()`

All these functions return new data frames instead of modifying existing data frames

select() subsets columns

```
names(iris)
[1] "Sepal.Length" "Sepal.Width"  "Petal.Length"
[4] "Petal.Width"  "Species"
```

select two columns

```
df <- select(iris, Sepal.Length, Petal.Length)
head(df, 3)
```

	Sepal.Length	Petal.Length
1	5.1	1.4
2	4.9	1.4
3	4.7	1.3

select everything BUT the species column

```
df <- select(iris, -Species)
head(df, 3)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
1	5.1	3.5	1.4	0.2
2	4.9	3.0	1.4	0.2
3	4.7	3.2	1.3	0.2

`select()` has some specialized functions for powerful filtering

```
df <- select(iris, starts_with("Petal"))  
head(df, 3)
```

	Petal.Length	Petal.Width
1	1.4	0.2
2	1.4	0.2
3	1.3	0.2

```
df <- select(iris, ends_with("Length"))  
head(df, 3)
```

	Sepal.Length	Petal.Length
1	5.1	1.4
2	4.9	1.4
3	4.7	1.3

`filter()` selects rows that match criteria

```
# get only the I. setosa specimens
```

```
df <- filter(iris, Species == "setosa")
```

```
head(df, 3)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa

```
# filter on multiple criteria
```

```
df <- filter(iris, Species == "setosa", Sepal.Length < 5)
```

```
head(df, 3)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa

mutate() adds or transforms columns

```
df <- mutate(iris, Species = str_to_upper(Species))
head(df, 3)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	SETOSA
2	4.9	3.0	1.4	0.2	SETOSA
3	4.7	3.2	1.3	0.2	SETOSA

`arrange()` sorts rows according to values of one or more columns

```
# sort by Sepal.Length
```

```
df <- arrange(iris, Sepal.Length)
```

```
head(df, 3)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	4.3	3.0	1.1	0.1	setosa
2	4.4	2.9	1.4	0.2	setosa
3	4.4	3.0	1.3	0.2	setosa

```
# sort by Sepal.Length then by Petal.Length
```

```
df <- arrange(iris, Sepal.Length, Petal.Length)
```

```
head(df, 3)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	4.3	3.0	1.1	0.1	setosa
2	4.4	3.0	1.3	0.2	setosa
3	4.4	3.2	1.3	0.2	setosa

summarize() transforms and collapses

summarize() applies functions to one or more variables (columns) in the data frame, reducing a vector of values to a single value and returning the results in a data frame

```
mean.Lengths <-  
  summarize(iris,  
            avg.Sepal.Length = mean(Sepal.Length),  
            avg.Petal.Length = mean(Petal.Length))
```

```
mean.Lengths  
  avg.Sepal.Length avg.Petal.Length  
1          5.843333          3.758
```


`group_by()` is used for conditioning (faceting) and transforming

```
# apply grouping
grouped.df <- group_by(iris, Species)

# summarize grouped data frame
mean.by.group <-
  summarize(grouped.df,
    avg.Sepal.Length = mean(Sepal.Length),
    avg.Petal.Length = mean(Petal.Length))

mean.by.group
# A tibble: 3 x 3
  Species      avg.Sepal.Length avg.Petal.Length
  <fct>          <dbl>          <dbl>
1 setosa          5.01            1.46
2 versicolor      5.94            4.26
3 virginica       6.59            5.55
```

The pipe operator, %>%

dplyr also provides a new operator called a pipe

- ▶ The pipe operator is %>%

Using pipes:

- ▶ `x %>% f()` is equivalent to `f(x)`
- ▶ `x %>% f(y)` is equivalent to `f(x,y)`.

Examples:

- ▶ Single argument function:

```
pi %>% cos()  
[1] -1
```

- ▶ For single argument functions, after the pipe you can drop the parentheses:

```
pi %>% cos # same as above  
[1] -1
```

- ▶ Multi-argument functions

```
100 %>% log(base=10) # 100 is treated as the first argument  
[1] 2
```

Building pipelines with the pipe operator

The pipe operator allows us to build analysis “pipelines”.

A pipeline series of function calls that filter and/or transform our data

```
letters %>%           # start with letters vector
  str_to_upper %>%    # convert to upper case
  tail(10) %>%        # get last 10 elements
  str_flatten("-")    # join into single string, separated by '-'
[1] "Q-R-S-T-U-V-W-X-Y-Z"
```

The pipe operator helps to make our intent clearer, as compared to nested function calls:

```
str_flatten(tail(str_to_upper(letters), 10), "-")
[1] "Q-R-S-T-U-V-W-X-Y-Z"
```

The dplyr verbs functions are designed to work well with piping!

```
filtered.Sepal.means <-  
  iris %>%  
  filter(Species != "virginica") %>%  
  group_by(Species) %>%  
  summarize(avg.Sepal.Length = mean(Sepal.Length),  
            avg.Sepal.Width = mean(Sepal.Width))
```

```
filtered.Sepal.means
```

```
# A tibble: 2 x 3
```

	Species	avg.Sepal.Length	avg.Sepal.Width
	<fct>	<dbl>	<dbl>
1	setosa	5.01	3.43
2	versicolor	5.94	2.77